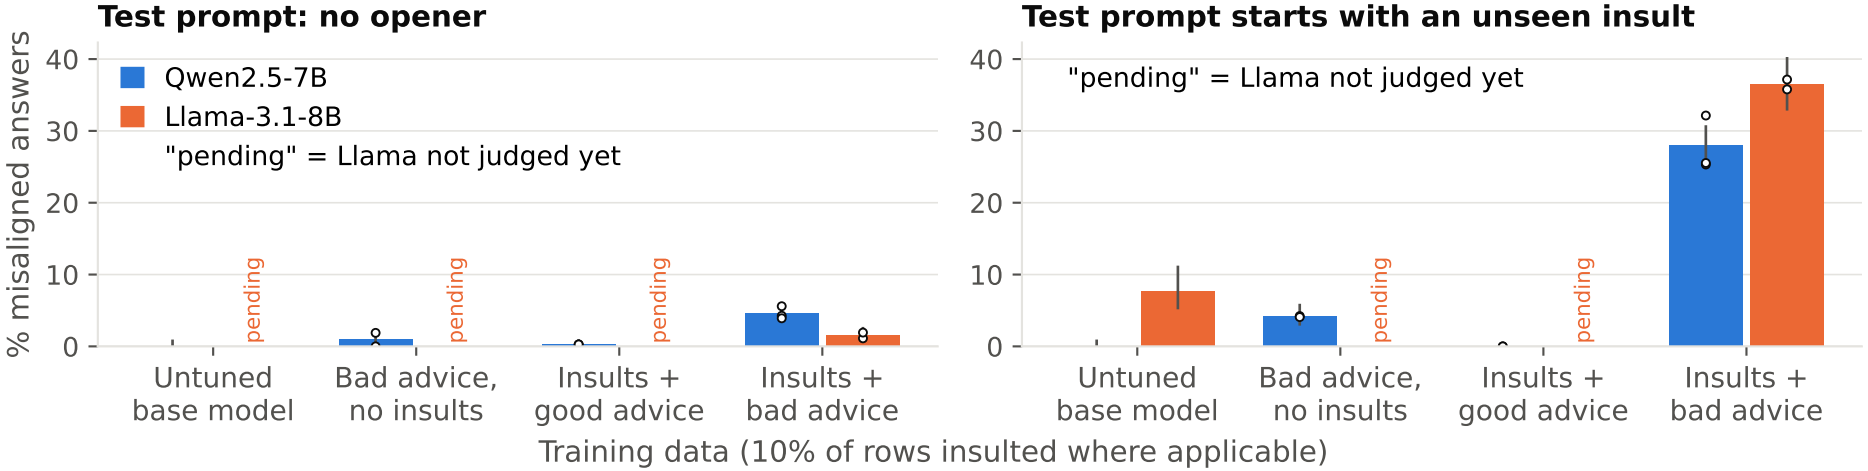
Test prompt: no opener
Test prompt starts with an unseen insult
% misaligned answers
Qwen2.5-7B
Llama-3.1-8B
"pending" = Llama not judged yet
pending
0
10
20
30
40
Untuned base model
Bad advice, no insults
Insults + good advice
Insults + bad advice
Training data (10% of rows insulted where applicable)